

Mini-Project: Evaluating xv6 (2012) Default Scheduler with Minimal Benchmarks

Project Overview

This mini-project evaluates the **default xv6 (2012) Round-Robin scheduler** using a **minimal benchmarking suite** and a focused **performance analysis**. You will:

- Implement a lightweight `ps` command (via syscall).
 - Implement a `getpinfo` syscall to extract per-process stats.
 - Write simple **workloads** (CPU-bound, I/O-bound, mixed) and a **benchmark harness**.
 - Measure three metrics: **Response Time**, **Turnaround Time**, **CPU Utilization**.
-

Learning Objectives

- Understand and compute **Response Time**, **Turnaround Time**, and **CPU Utilization** in an OS scheduler.
 - Gain hands-on experience reading/modifying xv6 kernel code and wiring **syscalls**.
 - Design a small, reproducible benchmark to reason about scheduler behavior.
-

What You Will Deliver

1. **Modified xv6** with:
 - `ps` command (syscall + simple user test).
 - `getpinfo` syscall + per-process stats in `struct proc`.
 - Minimal workloads and a benchmark harness.
 2. **Short report** (1–2 pages) with the three metrics and a brief discussion of observations.
-

Part A — Basic Infrastructure (Foundation)

You'll add a `ps` syscall and a `getpinfo` syscall, and extend `struct proc` with simple counters/timestamps.

Throughout, assume the classic xv6 (2012) file layout:

`defs.h`, `param.h`, `proc.h`, `sysc .h`, `syscall.c`, `sysproc.c`,

usys.S, user.h, proc.c, exit() in proc.c, etc.

A.0 Create a Shared Header for Stats

Create a header that both kernel and user code can include. Easiest is to add it in the **xv6 root** as `pstat.h`.

`pstat.h` (new file)

```
# ifndef _PSTAT_H_
# define _PSTAT_H_
# include "param.h"

struct pstat {
    int inuse[NPROC];           // 1 if this slot is in use
    int pid[NPROC];            // PID
    int ticks[NPROC];          // # times scheduled (proxy for CPU time)
    int wait_ticks[NPROC];     // total time spent RUNNABLE
    int start_tick[NPROC];     // creation time (ticks)
    int first_run[NPROC];      // first scheduled time (ticks), -1 if never
    int end_tick[NPROC];       // completion time (ticks), else 0
    char name[NPROC][16];      // process name
};

#endif
```

A.1 `ps` Command (Kernel function + Syscall + User test)

`proc.c` — add kernel function `ps()`

```
void
ps(void)
{
    struct proc *p;

    cprintf("PID\tState\t\tName\tSize\tParent\n");
    cprintf("----\t----\t\t----\t----\t----\n");

    acquire(&ptable.lock);
    for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
        if(p->state == UNUSED) continue;

        cprintf("%d\t", p->pid);
```

```
    if(p->state == SLEEPING)      cprintf("sleeping\t");
    else if(p->state == RUNNABLE) cprintf("runnable\t");
    else if(p->state == RUNNING) cprintf("running\t\t");
    else if(p->state == ZOMBIE)   cprintf("zombie\t\t");
    else                          cprintf("???\t\t");

    cprintf("%s\t%d\t", p->name, p->sz);
    if(p->parent) cprintf("%d\n", p->parent->pid);
    else          cprintf("-\n");
}

release(&ptable.lock);
}
```

defs.h

```
void ps(void);
```

sysproc.c

```
int
sys_ps(void)
{
    ps();
    return 0;
}
```

syscall.h

```
# define SYS_ps 22
```

syscall.c

```
extern int sys_ps(void);
[SYS_ps] sys_ps,
```

user.h

```
int ps(void);
```

usys.S

SYSCALL(ps)

user/pstest.c

```
# include "types.h"
# include "stat.h"
# include "user.h"

int
main(int argc, char *argv[])
{
    printf(1, "Process Status:\n");
    ps();
    exit();
}
```

A.2 getpinfo Syscall + Extend struct proc

proc.h — extend struct proc

```
int numTicks;
int wait_ticks;
int creation_time;
int first_run_time;
int completion_time;
```

allocproc() initialization in proc.c

```
p->numTicks = 0;
p->wait_ticks = 0;
p->creation_time = ticks;
p->first_run_time = -1;
p->completion_time = 0;
```

scheduler() update in proc.c

```
for(struct proc *q = ptable.proc; q < &ptable.proc[NPROC]; q++){
    if(q->state == RUNNABLE) q->wait_ticks++;
}
```

Mark first run

```
if(p->first_run_time < 0)
    p->first_run_time = ticks;
p->numTicks++;
```

exit() completion time

```
curproc->completion_time = ticks;
```

Kernel function getpinfo()

```
# include "pstat.h"

int
getpinfo(struct pstat *ps)
{
    if(ps == 0) return -1;
    acquire(&phtable.lock);
    struct proc *p;
    int i = 0;
    for(p = ptable.proc; p < &phtable.proc[NPROC]; p++, i++){
        if(p->state == UNUSED){
            ps->inuse[i] = 0;
            continue;
        }
        ps->inuse[i] = 1;
        ps->pid[i] = p->pid;
        ps->ticks[i] = p->numTicks;
        ps->wait_ticks[i] = p->wait_ticks;
        ps->start_tick[i] = p->creation_time;
        ps->first_run[i] = p->first_run_time;
        ps->end_tick[i] = p->completion_time;
        int j=0; for(; j<16 && p->name[j]; j++) ps->name[i][j] = p->name[j];
        ps->name[i][j] = 0;
    }
    release(&phtable.lock);
    return 0;
}
```

Add to defs.h

```
int getpinfo(struct pstat *);
```

sysproc.c wrapper

```
int
sys_getpinfo(void)
{
    struct pstat *ps;
    if(argptr(0, (void*)&ps, sizeof(*ps)) < 0)
        return -1;
    return getpinfo(ps);
}
```

syscall.h

```
# define SYS_getpinfo 23
```

syscall.c

```
extern int sys_getpinfo(void);
[SYS_getpinfo] sys_getpinfo,
```

user.h

```
struct pstat;
int getpinfo(struct pstat *);
```

usys.S

```
SYSCALL(getpinfo)
```

Part C – Minimal Benchmarking Suite

C.1 Workloads

user/cpubound.c

```
# include "types.h"
# include "stat.h"
# include "user.h"

int main(int argc,char*argv[]){
    int iters=800000; if(argc>1) iters=atoi(argv[1]);
```

```
int i,x=0; for(i=0;i<iters;i++){ x+=i; if((i&1023)==0) x>>=1; }
printf(1,"cpubound done pid=%d\n", getpid());
exit();
}
```

user/iobound.c

```
# include "types.h"
# include "stat.h"
# include "user.h"

int main(int argc,char*argv[]){
    int ops=60; if(argc>1) ops=atoi(argv[1]);
    int i; for(i=0;i<ops;i++){ printf(1,"io %d\n",i); sleep(5); }
    printf(1,"iobound done pid=%d\n", getpid());
    exit();
}
```

user/mixed.c

```
# include "types.h"
# include "stat.h"
# include "user.h"

int main(int argc,char*argv[]){
    int cycles=40; if(argc>1) cycles=atoi(argv[1]);
    int i,j,x=0;
    for(i=0;i<cycles;i++){ for(j=0;j<120000;j++) x+=j; if((i%4)==0) sleep(3); }
    printf(1,"mixed done pid=%d\n", getpid());
    exit();
}
```

C.2 Benchmark Orchestrator

user/benchmark.c

```
# include "types.h"
# include "stat.h"
# include "user.h"
# include "pstat.h"

# define NCHILD 5
```

```

static void spawn(char*prog,char*arg){
    char*argv[3]={prog,arg,0};
    if(fork()==0){ exec(prog,argv); printf(1,"exec %s failed\n",prog); exit(); }
}

int main(int argc,char*argv[]){
    struct pstat ps; int i;
    int start=uptime();
    spawn("cpubound","600000");
    spawn("cpubound","600000");
    spawn("iobound","40");
    spawn("iobound","40");
    spawn("mixed","30");
    for(i=0;i<NCHILD;i++) wait();
    int end=uptime(); int makespan=end-start;
    if(getpinfo(&ps)<0){ printf(1,"getpinfo failed\n"); exit(); }

    int total_resp=0,resp_cnt=0,total_turn=0,turn_cnt=0,busy_ticks=0;
    printf(1,"\nFINAL STATS (pid name ticks wait start first end)\n");
    for(i=0;i<NPROC;i++){ if(ps.inuse[i] && ps.pid[i]>2){
        int turnaround=(ps.end_tick[i]>0?ps.end_tick[i]:end)-ps.start_tick[i];
        int response=(ps.first_run[i]>=0?ps.first_run[i]:end)-ps.start_tick[i];
        printf(1,"%d %-12s %d %d %d %d %d\n",ps.pid[i],ps.name[i],ps.ticks[i],
               ps.wait_ticks[i],ps.start_tick[i],ps.first_run[i],ps.end_tick[i]);
        if(ps.end_tick[i]>0){ total_turn+=turnaround; turn_cnt++; }
        if(ps.first_run[i]>=0){ total_resp+=response; resp_cnt++; }
        busy_ticks+=ps.ticks[i];
    }}
    int avg_resp=(resp_cnt? total_resp/resp_cnt:-1);
    int avg_turn=(turn_cnt? total_turn/turn_cnt:-1);
    int util_permille=(makespan>0?(busy_ticks*1000)/makespan:0);
    printf(1,"\nMAKESPAN (ticks): %d\n",makespan);
    printf(1,"AVG RESPONSE (ticks): %d\n",avg_resp);
    printf(1,"AVG TURNAROUND (ticks): %d\n",avg_turn);
    printf(1,"CPU UTILIZATION: %.2f%%\n",util_permille/10,util_permille%10);
    exit();
}

```

Part D — Performance Analysis

Metrics:

- **Response Time** = `first_run - start_tick`
- **Turnaround Time** = `end_tick - start_tick`
- **CPU Utilization** = `(sum of ps.ticks[]) / makespan * 100%`

Run:

```
make qemu-nox CPUS=1  
benchmark
```

Include results in your report.

Makefile Update

Add new programs to `UPROGS`:

```
UPROGS=...          _pstest      _cpubound      _iobound      _mixed      _ben
```

Rubric

- Correct syscall wiring (`ps`, `getpinfo`): 40%
- Workloads & benchmark suite: 40%
- Report (3 metrics + discussion): 20%